

# Applied Data Mining: Homework #7

Keith Hickman

Due on December 6, 2017

## Problem 1

In this problem, you are asked to use SVM to predict whether a given car gets high or low gas mileage based on the Auto data set. The data set can be obtained as follows:

```
##install.packages("ISLR")
library(ISLR)
View(Auto)
```

### 1.1

Create a binary variable that takes on a 1 for cars with gas mileage above the median, and a 0 for cars with gas mileage below the median. Add this variable to the data as a new variable and name it as "mpglevel" (mpglevel is the response variable for questions 1.2 and 1.3).

### R Code

```
Auto$mpglevel <- as.factor(Auto$mpg >= median(Auto$mpg))
print(Auto$mpglevel)

## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [386] TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## Levels: FALSE TRUE

summary(Auto)

##      mpg      cylinders  displacement  horsepower
## Min.   : 9.00   Min.    :3.000   Min.    : 68.0   Min.    : 46.0
## 1st Qu.:17.00   1st Qu.:4.000   1st Qu.:105.0   1st Qu.: 75.0
## Median :22.75   Median :4.000   Median :151.0   Median : 93.5
## Mean   :23.45   Mean    :5.472   Mean    :194.4   Mean    :104.5
## 3rd Qu.:29.00   3rd Qu.:8.000   3rd Qu.:275.8   3rd Qu.:126.0
## Max.    :46.60   Max.    :8.000   Max.    :455.0   Max.    :230.0
##
##      weight  acceleration      year      origin
## Min.    :1613   Min.    : 8.00   Min.    :70.00   Min.    :1.000
## 1st Qu.:2225   1st Qu.:13.78   1st Qu.:73.00   1st Qu.:1.000
## Median :2804   Median :15.50   Median :76.00   Median :1.000
## Mean    :2978   Mean     :15.54   Mean     :75.98   Mean     :1.577
```

```
## 3rd Qu.:3615    3rd Qu.:17.02    3rd Qu.:79.00    3rd Qu.:2.000
## Max.      :5140    Max.      :24.80    Max.      :82.00    Max.      :3.000
##
##              name      mpglevel
## amc matador      : 5    FALSE:196
## ford pinto       : 5    TRUE :196
## toyota corolla   : 5
## amc gremlin      : 4
## amc hornet       : 4
## chevrolet chevette: 4
## (Other)         :365
```

## 1.2

Fit a linear support vector classifier to the data with various values of cost (cost = c(0.01, 0.1, 1, 5, 10, 100)), in order to predict whether a car gets high or low gas mileage. Report the cross-validation errors associated with different values of this parameter. Comment on your results, i.e., what is the cost value for the model that has the lowest cross-validation error?

### R Code

```
library(e1071)
set.seed(123456)
rndSample <- sample(1:nrow(Auto), 300)
tr <- Auto[rndSample, ]
ts <- Auto[-rndSample, ]

# the default svm () uses radial kernel with constraints violations of cost of 1
## ??svm

#Beginning with a cost of .01
s.01 <- svm(mpglevel ~ ., tr, C=.01)
ps.01 <- predict(s.01, ts)
cm.01 <- table(ps.01, ts$mpglevel) #confusion matrix
cm.01

##
## ps.01    FALSE TRUE
## FALSE    35    1
## TRUE     3    53

100*(1-sum(diag(cm.01))/sum(cm.01))

## [1] 4.347826

#Cost of .1
s.1 <- svm(mpglevel ~ ., tr, C=.1)
ps.1 <- predict(s.1, ts)
```

```

cm.1 <- table(ps.1, ts$mpglevel) #confusion matrix
100*(1-sum(diag(cm.1))/sum(cm.1))

## [1] 4.347826

#Default cost of 1
s1 <- svm(mpglevel ~ ., tr,C=1)
ps1 <- predict(s1, ts)
cm1 <- table(ps1, ts$mpglevel) #confusion matrix
100*(1-sum(diag(cm1))/sum(cm1))

## [1] 4.347826

##Cost of 5
s5 <- svm(mpglevel ~ ., tr,C=5)
ps5 <- predict(s5, ts)
cm5 <- table(ps5, ts$mpglevel) #confusion matrix
100*(1-sum(diag(cm5))/sum(cm5))

## [1] 4.347826

## Cost of 10
s10 <- svm(mpglevel ~ ., tr,C=10)
ps10 <- predict(s10, ts)
cm10 <- table(ps10, ts$mpglevel) #confusion matrix
100*(1-sum(diag(cm10))/sum(cm10))

## [1] 4.347826

## Cost of 100
s100 <- svm(mpglevel ~ ., tr,C=100)
ps100 <- predict(s100, ts)
cm100 <- table(ps100, ts$mpglevel) #confusion matrix
100*(1-sum(diag(cm100))/sum(cm100))

## [1] 4.347826

```

## Cross-validation Errors and Discussion of the Results

All of my cross-validation errors are the same with the costs from .01 to 100 = 6.52137% error rate.

### 1.3

Now repeat (1.2), this time using SVMs with radial and polynomial basis kernels, with different values of gamma (c(0.01, 0.1, 1, 5, 10, 100)) and degree (c(2, 3, 4)) and cost (c(0.1, 1, 5, 10)). Use the cost and degree parameters values for polynomial kernels. The cost and gamma parameters values are given for radial basis kernels. Comment on your results, i.e., what are the parameters values (cost, degree, gamma) for the model that has the lowest cross-validation error?

## R Code

```
#Low Cost, Gamma, and Degree
svm1 <- svm(mpglevel ~ ., tr,C=1, degree=1, gamma=1)
ps111 <- predict(svm1, ts)
cm111 <- table(ps111, ts$mpglevel) #confusion matrix
100*(1-sum(diag(cm111))/sum(cm111))

## [1] 3.26087

#High Cost, Gamma, Degree
s100 <- svm(mpglevel ~ ., tr,C=100, degree=3, gamma=10)
ps100 <- predict(s100, ts)
cm100 <- table(ps100, ts$mpglevel) #confusion matrix
100*(1-sum(diag(cm100))/sum(cm100))

## [1] 55.43478
```

## Discussion of Results

I find that modifying the cost doesn't change the overall CV error rate, but that lower gamma and degree parameters has a significant impact on the CV error rate.

## Problem 2

```
##install.packages("dplyr")
##library(dplyr)
View(Caravan)
```

### 2.1

Create a training set consisting of the first 1,000 observations, and a test set consisting of the remaining observations. The class variable is "Purchase" whose values are "No" and "Yes". Transform "No" to 0 "Yes" to 1. Place the R code below.

## R Code

```
Caravan$Purchase <- as.character(ifelse(Caravan$Purchase=="Yes", 1, 0))
train <- Caravan[1:1000,]
test <- Caravan[1001:5822,]
## View(Caravan)
typeof(Caravan$Purchase)

## [1] "character"
```

I kept getting "Nan" values when evaluating my gbm model summary. I tried converting the Purchase variable to a factor, character, and integer. After looking through the text and lecture notes, I went to stack overflow, but the suggestions there didn't help either.

## 2.2

Fit a boosting model to the training set with Purchase as the response and the other variables as predictors. Use 1,000 trees, and a shrinkage value of 0.01. Which predictors appear to be the most important?

### R Code

```
##install.packages("gbm")
library(gbm)

## Loading required package: survival
## Loading required package: lattice
## Loading required package: splines
## Loading required package: parallel
## Loaded gbm 2.1.3

model <- gbm(Purchase ~ ., data=test, n.trees = 1000, shrinkage = .01)

## Distribution not specified, assuming bernoulli ...

summary(model, plotit = FALSE)

##           var      rel.inf
## PPERSAUT PPERSAUT 25.07662675
## PPLEZIER PPLEZIER 14.08425902
## PBRAND   PBRAND  11.16089257
## MOPLLAAG MOPLLAAG  5.21533664
## MINKGEM  MINKGEM  4.69512214
## ALEVEN   ALEVEN  4.18820749
## APERSAUT APERSAUT  3.25526666
## PBYSTAND PBYSTAND  3.07208210
## MBERMIDD MBERMIDD  2.30144136
## MOSTYPE  MOSTYPE  2.20101843
## MBERHOOG MBERHOOG  1.81425545
## MBERARBG MBERARBG  1.68668522
## MAUT1    MAUT1    1.58667411
## MKOOPKLA MKOOPKLA  1.52103448
## PWAPART  PWAPART  1.45513613
## MGODOV   MGODOV   1.29436395
## MINK7512 MINK7512  1.29145049
## AFIETS   AFIETS   1.26061847
## MINKM30  MINKM30  1.05005438
## PGEZONG  PGEZONG  1.03874797
## MSKC     MSKC     1.00708097
## PFIETS   PFIETS   0.98458131
## MOSHOOFD MOSHOOFD  0.96505929
## MOPLMIDD MOPLMIDD  0.81288729
```

##	MINK3045	MINK3045	0.70162814
##	MOPLHOOG	MOPLHOOG	0.62782018
##	MGODGE	MGODGE	0.56303116
##	MRELGE	MRELGE	0.53573276
##	MHHUUR	MHHUUR	0.53067741
##	MSKA	MSKA	0.50750504
##	PLEVEN	PLEVEN	0.40247868
##	MGODPR	MGODPR	0.39419935
##	MINK4575	MINK4575	0.35472250
##	MAUT0	MAUT0	0.23626821
##	MBERBOER	MBERBOER	0.21297093
##	MINK123M	MINK123M	0.18901388
##	MFWEKIND	MFWEKIND	0.18129730
##	MSKB1	MSKB1	0.17605477
##	MGODRK	MGODRK	0.17445922
##	MZPART	MZPART	0.17139814
##	MRELSA	MRELSA	0.14721546
##	MHKOOP	MHKOOP	0.14248855
##	MSKD	MSKD	0.13100341
##	MZFONDS	MZFONDS	0.11782199
##	PINBOED	PINBOED	0.11175351
##	MBERZELF	MBERZELF	0.10728261
##	MFGEKIND	MFGEKIND	0.09740308
##	MGEMLEEF	MGEMLEEF	0.07722697
##	MAUT2	MAUT2	0.05428475
##	PWALAND	PWALAND	0.03537936
##	MAANTHUI	MAANTHUI	0.00000000
##	MGEMOMV	MGEMOMV	0.00000000
##	MRELOV	MRELOV	0.00000000
##	MFALLEEN	MFALLEEN	0.00000000
##	MBERARBO	MBERARBO	0.00000000
##	MSKB2	MSKB2	0.00000000
##	PWABEDR	PWABEDR	0.00000000
##	PBESAUT	PBESAUT	0.00000000
##	PMOTSCO	PMOTSCO	0.00000000
##	PVRAAUT	PVRAAUT	0.00000000
##	PAANHANG	PAANHANG	0.00000000
##	PTRACTOR	PTRACTOR	0.00000000
##	PWERKT	PWERKT	0.00000000
##	PBROM	PBROM	0.00000000
##	PPERSONG	PPERSONG	0.00000000
##	PWAOREG	PWAOREG	0.00000000
##	PZEILPL	PZEILPL	0.00000000
##	AWAPART	AWAPART	0.00000000
##	AWABEDR	AWABEDR	0.00000000
##	AWALAND	AWALAND	0.00000000
##	ABESAUT	ABESAUT	0.00000000
##	AMOTSCO	AMOTSCO	0.00000000
##	AVRAAUT	AVRAAUT	0.00000000
##	AAANHANG	AAANHANG	0.00000000

```
## ATTRACTOR ATTRACTOR 0.00000000
## AWERKT AWERKT 0.00000000
## ABROM ABROM 0.00000000
## APERSONG APERSONG 0.00000000
## AGEZONG AGEZONG 0.00000000
## AWAOREG AWAOREG 0.00000000
## ABRAND ABRAND 0.00000000
## AZEILPL AZEILPL 0.00000000
## APLEZIER APLEZIER 0.00000000
## AINBOED AINBOED 0.00000000
## ABYSTAND ABYSTAND 0.00000000
```

### Problem 3

```
library(data.table)
library("curl")
mydata <- fread("https://archive.ics.uci.edu/ml/machine-learning-databases/ionosphere/ionosphere.data")
mydata <- as.data.frame(mydata)
mydata <- mydata[, -2] #remove the second variable
```

#### 3.1

Create a training data set containing a random sample of 300 data points and a test set containing the remaining observations. Name the training data and test data as mydata.training and mydata.testing, respectively. Place the R code below. You will use mydata.training and mydata.testing to answer rest of the questions. Thus, create them once and use mydata.training to train the models (classifiers) and mydata.testing to test the models. The last variable variable (35th variable in the data) is the response and the other variables are predictors.

#### R Code

```
set.seed(1234)
rndSample <- sample(1:nrow(mydata), 300)
mydata.training <- mydata[rndSample,]
mydata.testing <- mydata[-rndSample,]
```

#### 3.2

Train a naive bayes classifier using 10-fold cross-validation over mydata.training. Use this model to predict the observations in mydata.testing. Form a confusion matrix and report the error rate of the classifier over mydata.testing.

```
##install.packages("lme4", dependencies = TRUE)
##library(lme4)
##methods(sigma)
##install.packages("pbkrtest", dependencies = TRUE)
##install.packages("DEoptimR")
##install.packages("caret", dependencies = TRUE)
```

```

## library(caret)
##library(e1071)
head(mydata)

##      V1      V3      V4      V5      V6      V7      V8      V9      V10
## 1  1 0.99539 -0.05889  0.85243  0.02306  0.83398 -0.37708  1.00000  0.03760
## 2  1 1.00000 -0.18829  0.93035 -0.36156 -0.10868 -0.93597  1.00000 -0.04549
## 3  1 1.00000 -0.03365  1.00000  0.00485  1.00000 -0.12062  0.88965  0.01198
## 4  1 1.00000 -0.45161  1.00000  1.00000  0.71216 -1.00000  0.00000  0.00000
## 5  1 1.00000 -0.02401  0.94140  0.06531  0.92106 -0.23255  0.77152 -0.16399
## 6  1 0.02337 -0.00592 -0.09924 -0.11949 -0.00763 -0.11824  0.14706  0.06637
##      V11      V12      V13      V14      V15      V16      V17      V18
## 1 0.85243 -0.17755  0.59755 -0.44945  0.60536 -0.38223  0.84356 -0.38542
## 2 0.50874 -0.67743  0.34432 -0.69707 -0.51685 -0.97515  0.05499 -0.62237
## 3 0.73082  0.05346  0.85443  0.00827  0.54591  0.00299  0.83775 -0.13644
## 4 0.00000  0.00000  0.00000  0.00000 -1.00000  0.14516  0.54094 -0.39330
## 5 0.52798 -0.20275  0.56409 -0.00712  0.34395 -0.27457  0.52940 -0.21780
## 6 0.03786 -0.06302  0.00000  0.00000 -0.04572 -0.15540 -0.00343 -0.10196
##      V19      V20      V21      V22      V23      V24      V25      V26
## 1 0.58212 -0.32192  0.56971 -0.29674  0.36946 -0.47357  0.56811 -0.51171
## 2 0.33109 -1.00000 -0.13151 -0.45300 -0.18056 -0.35734 -0.20332 -0.26569
## 3 0.75535 -0.08540  0.70887 -0.27502  0.43385 -0.12062  0.57528 -0.40220
## 4 -1.00000 -0.54467 -0.69975  1.00000  0.00000  0.00000  1.00000  0.90695
## 5 0.45107 -0.17813  0.05982 -0.35575  0.02309 -0.52879  0.03286 -0.65158
## 6 -0.11575 -0.05414  0.01838  0.03669  0.01519  0.00888  0.03513 -0.01535
##      V27      V28      V29      V30      V31      V32      V33      V34
## 1 0.41078 -0.46168  0.21266 -0.34090  0.42267 -0.54487  0.18641 -0.45300
## 2 -0.20468 -0.18401 -0.19040 -0.11593 -0.16626 -0.06288 -0.13738 -0.02447
## 3 0.58984 -0.22145  0.43100 -0.17365  0.60436 -0.24180  0.56045 -0.38238
## 4 0.51613  1.00000  1.00000 -0.20099  0.25682  1.00000 -0.32382  1.00000
## 5 0.13290 -0.53206  0.02431 -0.62197 -0.05707 -0.59573 -0.04608 -0.65697
## 6 -0.03240  0.09223 -0.07859  0.00732  0.00000  0.00000 -0.00039  0.12011
##      V35
## 1  g
## 2  b
## 3  g
## 4  b
## 5  g
## 6  b

##model = train(mydata.training, 'nb', trControl=trainControl(method='cv', number=10))
##model <- NaiveBayes(mydata.training$ ~ ., data = tr)
##predict(model, mydata.testing)
##table(predict(model, mydata.testing))
##plot(model)

```

Unfortunately, I could not load the library caret. I tried several fixes, including uninstalling and reinstalling the source files, updating R (current version is 3.3.3), installing several



other packages ahead of the caret, and it continually gives a namespace error and will not install. Therefore I can't use the `train()` function.

I am going to re-attempt fixes tomorrow.